OpenStreetMap Data Case Study - Project 3

Author: Jaks Wildman

Map Area: I selected Columbus, Ohio, USA as it is the city I've spent the most time in.

- Location: Columbus, Ohio
- OpenStreetMap URL

Data Audit

Unique Tags

The XML file utilizes many tags to structure the data. Using `mapparser.py` I counted the number of each unique tag respectively from columbusOH.osm. The code for this is `mapparser.py` taken from the class case study.

- `'node' : 1502751`
- `'member' : 38206,`
- `'nd' : 1837440,`
- `'tag' : 687442,`
- `'bounds' : 1,` S
- `'note' : 526,`
- `'meta' : 91356,`
- `'relation' : 77866`
- `'way' : 177854`
- `'osm' : 1`

Patterns

This set of pattern checks was run on the entire osm file vs the sample that's in the root folder and utilizes regular expressions. `tags.py` contains the necessary code to count these 4 categories.

- `"lower" : 429525` - Tags that only contain lowercase letters and pass the validity checks.
- `"lower_colon" : 244778` - Tags that are otherwise valid, yet have colons in their name.
- `"problemchars" : 0` - Tags with problematic characters such as "=", "+", "&", "," , "?" and more.
- `"other" : 13139` - Any other tags that don't fit in the 3 prior categories

Problems In Data

The most significant issue that had to be addressed was the inconsistency in the street names and their abbreviations.
In order to correct these the following functions were utilized within the `audit.py` file:

- `audit_street_type` : Determines if the street name is within the list of expected names
- `street_name` : Tests whether the 'k' attribute matches the key for street data (`addr:street`)
- `audit` : Returns a dictionary of key, value pairs which meet the criteria in the preceding functions
- `update_name` : Actually does the update on the street name

Data Overview

File Sizes

- `columbusOH.osm` : 324 MB
- `nodes.csv` : 124MB
- `nodes_tags.csv` : 3.37 MB
- `ways.csv` : 10.5 MB
- `ways_nodes.csv` : 43.7 MB
- `ways_tags.csv` : 20.2 MB
- `cbus.db` : 171 MB

Number of Nodes:

```
sqlite> SELECT COUNT(*) FROM NODES
```

Output: `1502751`

Number of Ways:

```
sqlite> SELECT COUNT(*) FROM NODES
```

Output: `177854`

A Count by type of the top 15 Node Tags

```
sqlite> SELECT DISTINCT TYPE, COUNT(ID) as TYPE_COUNT
 FROM NODES_TAGS
 GROUP BY TYPE
 ORDER BY TYPE_COUNT DESC
 LIMIT 15 ;
```

Output:

```
regular         | 76620
addr            | 7879
gnis            | 4409
species         | 1396
fire_hydrant    | 572
traffic_signals | 404
brand           | 340
contact         | 121
payment         | 79
tower           | 77
name            | 50
xmas            | 46
historic        | 42
service         | 39
surveillance    | 38
```

Number of Unique users:

```
sqlite> SELECT COUNT(DISTINCT(u.uid))
 FROM (SELECT uid FROM nodes
       UNION ALL
       SELECT uid from ways) u;
```

Output: `1151`

Top Contributors:

```
sqlite> SELECT USER, COUNT(*) AS EDITS
  FROM (SELECT USER FROM NODES
        UNION ALL
        SELECT USER FROM WAYS)
  GROUP BY USER
  ORDER BY EDITS DESC
  LIMIT 10;
```

Output:

```
woodpeck_fixbot  | 211799
doktorpixel14    | 167004
Anonononon       | 157726
Nimbalo          | 150409
MerlinPendragon  | 108529
duck57           | 88834
AndrewSP37       | 87877
Vid the Kid      | 69057
kbzimmer         | 61976
St-Motel         | 53533
```

Popular Restaurants by Cuisine

```
sqlite> SELECT NODES_TAGS.VALUE, COUNT(*) AS NUM
  FROM NODES_TAGS
  JOIN (SELECT DISTINCT(ID) FROM NODES_TAGS
        WHERE VALUE="restaurant") r on nodes_tags.id=r.id
  WHERE NODES_TAGS.KEY = 'cuisine'
```

```
    GROUP BY nodes_tags.value
    ORDER BY NUM DESC
    LIMIT 15;
```

Output:

```
pizza           | 33
american        | 26
chinese         | 22
mexican         | 16
sandwich        | 14
italian         | 10
asian           | 8
ice_cream       | 8
indian          | 6
greek           | 5
japanese        | 5
sushi           | 5
burger          | 4
chicken;american | 3
barbecue        | 2
```

Greatest number of Worship Centers by Religion:

```
sqlite> SELECT NODES_TAGS.VALUE, COUNT(*) AS NUM
  FROM NODES_TAGS
  JOIN (SELECT DISTINCT(ID) FROM NODES_TAGS
        WHERE VALUE = "place_of_worship") a
    ON NODES_TAGS.ID = A.ID
  WHERE NODES_TAGS.KEY = "religion"
  GROUP BY NODES_TAGS.VALUE
  ORDER BY NUM DESC
  LIMIT 3;
```

Output:

```
christian   | 536
muslim      | 2
jewish      | 1
```

Other Cities in the data by count

```
sqlite>SELECT tags.value, COUNT(*) as count
  FROM (SELECT * FROM nodes_tags UNION ALL
        SELECT * FROM ways_tags) tags
 WHERE tags.key LIKE '%city'
 GROUP BY tags.value
 ORDER BY count DESC
 limit 15;
```

Output:

```
Columbus        | 3215
Upper Arlington | 492
Dublin          | 403
Gahanna         | 170
Westerville     | 150
Lockbourne      | 130
Hilliard        | 95
Worthington     | 65
Pickerington    | 64
15              | 62
Pataskala       | 55
Galloway        | 33
Reynoldsburg    | 33
Grove City      | 32
6               | 27
```

Not only is this NOT just Columbus, but it confirms that there would need to be more cleaning efforts to identify what should really be in place of the numerical "city" values.

Conclusion

Considering the amount of data, and how it is all manually added it is surprisingly clean. However that doesn't mean it IS clean. Each layer reveals more issues that would need to be addressed such as returning lists of counties. Many times there are more than one in the result. Due to the size of the data, and the non-standard tags it can be very easy to have duplication based on different colloquial names.

Suggestion: Set a reasonable standard on tags for usage so that things like restaurant:cuisine = american and chicken;american doesn't happen. I'm not quite sure what "chicken;america" means and based on it's usage it doesn't appear many others do either.

A particular constraint however is that standards are difficult to get wide spread use of and would also require retroactive cleaning. It's very possible that it could simply make it worse or be cost prohibitive to implement.

A parser could be implemented to control the input as well, however that will slow the ability to do bulk uploads as they would need to be changed for the parser.

Files

- Case_study/ : All the scripts used and modified in the OSM Case study
- README.md : This File
- columbusOH_sample.osm: sampling of ColumbusOH.osm, taking every 5th element.
- audit.py: Performs the audit of the street and updates the names
- data.py: Creates the csvs, parses and shapes them
- create_db.py: Creates the database and fills it with the data from the csvs
- mapparser.py: Find the unique tags within the dataset
- osm_sampling.py: creates the osm_sample file
- tags.py: counts multiple patterns in the Tags
- users.py: counts unique users
- P3-OSM-Data Wrangling with SQL.pdf: The report

Sources

- https://stackoverflow.com/questions/2392732/sqlite-python-unicode-and-non-utf-data
- http://puwenning.github.io/2016/02/10/P3-project-openstreetmap-data-case-study/